




ctbench: compile time benchmarking for Clang

Jules Pénuchot



CODE RECKONS

Science to the CORE

ctbench

compile time benchmarking for Clang

Jules Pénuchoth | he/him

PhD Student

Parallel Systems, LISN, Paris-Saclay University

ctbench

compile time benchmarking for Clang

ctbench

set of tools for compile time benchmarking and analysis for Clang

ctbench

compile time benchmarking for Clang

Metaprogramming is evolving

- Support libraries

- Boost.Mpl,
- Boost.Fusion,
- Boost.Hana,
- Boost.Mp11,
- Brigand,
- *that code snippet repository you probably own*

→ People want and build better metaprogramming libraries as time goes

ctbench

compile time benchmarking for Clang

Metaprogramming is evolving

- C++ itself
 - Type traits (C++11)
 - constexpr specifier (C++11)
 - if constexpr (C++17)
 - More constexpr containers (C++20) - <https://wg21.link/p0784>
 - Reflection proposal - <https://wg21.link/p1974>
 - Metaprogramming proposal - <https://wg21.link/p2237>
 - ~~C++ will become self aware and take over the Metaverse~~

→ C++ might finally get a proper **metaprogramming API**

ctbench

compile time benchmarking for Clang

Metaprogramming is evolving

- Applications

- Eigen - <https://eigen.tuxfamily.org/>
- Blaze - <https://bitbucket.org/blaze-lib/blaze>
- CTRE - <https://github.com/hanickadot/compile-time-regular-expressions>
- CTPG - <https://github.com/peter-winter/ctpg>

→ Going towards... Compile time compilers?

ctbench

compile time benchmarking for Clang

Is it a good thing? Sure!

- Better **embedded domain-specific languages**
→ Makes C++ libraries easier to interact with for non C++-savvy audiences
- Metaprograms that **scale** better
→ Better performing tools and more regular project management
- Makes metaprogramming **mainstream**
→ Brings more people to the table for contributing to amazing projects

However...

ctbench

compile time benchmarking for Clang

<< With great constexpr power comes great **compile time**. >>

Or not..?

ctbench

compile time benchmarking for Clang

Metaprogramming lacks tooling

Metaprogramming is *almost* on par with regular programming...

- ...but regular programming has debuggers, profilers,
- we know how to benchmark it to get **meaningful**, quantitative results,
- no such process for metaprograms,
- little to no **science** behind compile time rule of thumbs.

→ We need a sane process for understanding compile times.

ctbench

compile time benchmarking for Clang

How to measure compile times?

- Templight, *Zoltán Borók-Nagy, Zoltán Porkoláb, and József Mihalicza* (2009)
- Metabench, *Louis Dionne and Bruno Dutra* (2016)
- Build-Bench, *Fred Tingaud* (2017)
- Clang time-trace & Clang Build Analyzer, *Aras Pranckevičius* (2019)

ctbench

compile time benchmarking for Clang

Templight, Zoltán Borók-Nagy, Zoltán Porkoláb, and József Mihalicza (2009)

<https://github.com/mikael-s-persson/templight>

- Clang-based profiling tool, used as a drop-in replacement,
- profiles time and memory usage of **template instantiations**,
- interactive debugging,
- mostly a one-shot profiler,
- templight-tools provide several conversion options (Graphviz, XML, etc.)

→ Strong effort but no backend data

ctbench

compile time benchmarking for Clang

Metabench, *Louis Dionne and Bruno Dutra* (2016)

<https://github.com/ldionne/metabench/>

- Framework + compile time benchmarks for **libraries**,
- generates benchmarks of different sizes using Ruby templates (ERB),
- **black box** approach (*only measures compiler execution time*),
- supports **all** compilers,
- web export

→ Valuable data, but harder to reuse

ctbench

compile time benchmarking for Clang

Build-Bench, *Fred Tingaud* (2017)

<https://build-bench.com/>

- Simple web interface,
- **black box** approach,
- supports **all** compilers,
- makes it easy to compare one-shot benchmarks,
- makes **per-compiler** comparisons possible

→ Nice to play with, but not extensive enough

ctbench

compile time benchmarking for Clang

Clang time-trace, *Aras Pranckevičius* (2019)

<https://github.com/aras-p/llvm-project-20170507/pull/2>

- **Code-aware profiler** provided by Clang,
- outputs flame graph **JSON** files for chrome://tracing,
- breaks down compile passes for **every symbol**,
- works for the frontend **and** the backend

→ Very strong reuse potential, available everywhere

ctbench

compile time benchmarking for Clang

Clang Build Analyzer, *Aras Pranckevičius* (2019)

<https://github.com/aras-p/ClangBuildAnalyzer>

- Built on-top of **time-trace**
- Provides a user-friendly output

→ Extracts essential data from a time-trace file

ctbench

compile time benchmarking for Clang

What is ctbench?

<https://github.com/jpenuchot/ctbench>

- Compile time benchmarking & **data analysis** tool for Clang,
- built on-top of **time-trace**,
- **repeatability** & accuracy in mind,
- **variable size** benchmarks,
- C++ developer friendly:
 - C++ only benchmark files,
 - CMake API,
 - JSON config files (with a few ones already provided)

But how does it work?

ctbench

compile time benchmarking for Clang

Benchmarking methodology

- Benchmark set:
 - collection of benchmark cases to compare
- Benchmark case:
 - compilable C++ file,
 - compiled several times for a given **range of sizes**,
 - benchmark **iteration size** passed as a preprocessor define
- Benchmark iteration:
 - terminology for a benchmark case compiled with a **given size**,
 - **several samples** for each iteration size for improved accuracy
- Sample:
 - **one** time-trace file

→ Benchmark set → Benchmark cases → Benchmark iterations → Samples 17/32

ctbench

compile time benchmarking for Clang

Benchmark case - Recursive sum

```
// Metaprogram to benchmark:
template <unsigned N> struct ct_uint_t { static constexpr unsigned value = N; };

template <typename T> auto sum(T const &) { return T::value; }
template <typename T, typename... Ts> auto sum(T const &, Ts const &...tl) {
    return T::value + sum(tl...);
}

// Benchmark driver:
#include <boost/preprocessor/repetition/enum.hpp>
#define GEN_MACRO(Z, N, TEXT) TEXT<N> {}
unsigned foo() {
    // return sum(ct_uint_t<1>{}, ..., ct_uint_t<BENCHMARK_SIZE>{});
    return sum(BOOST_PP_ENUM(BENCHMARK_SIZE, GEN_MACRO, ct_uint_t));
}
```

ctbench

compile time benchmarking for Clang

Benchmark case - Expansion sum

```
// Metaprogram to benchmark:
template <unsigned N> struct ct_uint_t { static constexpr unsigned value = N; };

template <typename... Ts> auto sum(Ts const &...) { return (Ts::value + ...); }

// Benchmark driver:
#include <boost/preprocessor/repetition/enum.hpp>
#define GEN_MACRO(Z, N, TEXT) TEXT<N> {}
unsigned foo() {
    // return sum(ct_uint_t<1>{}, ..., ct_uint_t<BENCHMARK_SIZE>{});
    return sum(BOOST_PP_ENUM(BENCHMARK_SIZE, GEN_MACRO, ct_uint_t));
}
```

→ *ctbench* defines *BENCHMARK_SIZE* for each iteration size

ctbench

compile time benchmarking for Clang

CMake API

- Benchmark declaration

```
ctbench_add_benchmark(variadic_sum.recursive # Target name
  variadic_sum/recursive.cpp                # Benchmark file
  1 32 1                                     # Start, stop, and step
  10)                                         # Number of repetitions
```

- Graph declaration

```
ctbench_add_graph(variadic_sum-graph # Target name
  configs/feature_comparison.json    # Config file
  variadic_sum.expansion              # Benchmark target
  variadic_sum.recursive)            # ...
```

- Optional: Bring your own flags with `ctbench_add_custom_benchmark`

ctbench

compile time benchmarking for Clang

JSON configs

- Graph options:
 - `plotter` (currently: `debug`, `simple curves`, and `stacked curves`),
 - output file format (Gnuplot based thanks to `Sciplot`),
 - graph dimensions,
 - other `plotter`-specific graph options.
- Data options:
 - define named data groups,
 - implement predicates to target specific data,
 - predicates: regex matching, JSON inclusion matching, logical operators

→ Pre-made configurations available

ctbench

compile time benchmarking for Clang

Demo time!

ctbench

compile time benchmarking for Clang

Rule of Chiel

Ordering of compile time impact for several meta-programming techniques:

Expensive	SFINAE
	Instantiating a function template
	Instantiating a type
	Calling an alias
	Adding a parameter to a type
	Adding a parameter to an alias call
Cheap	Looking up a memoized type

Source: "Type Based Template Metaprogramming is Not Dead", *Odin Holmes*
(C++Now 2017)

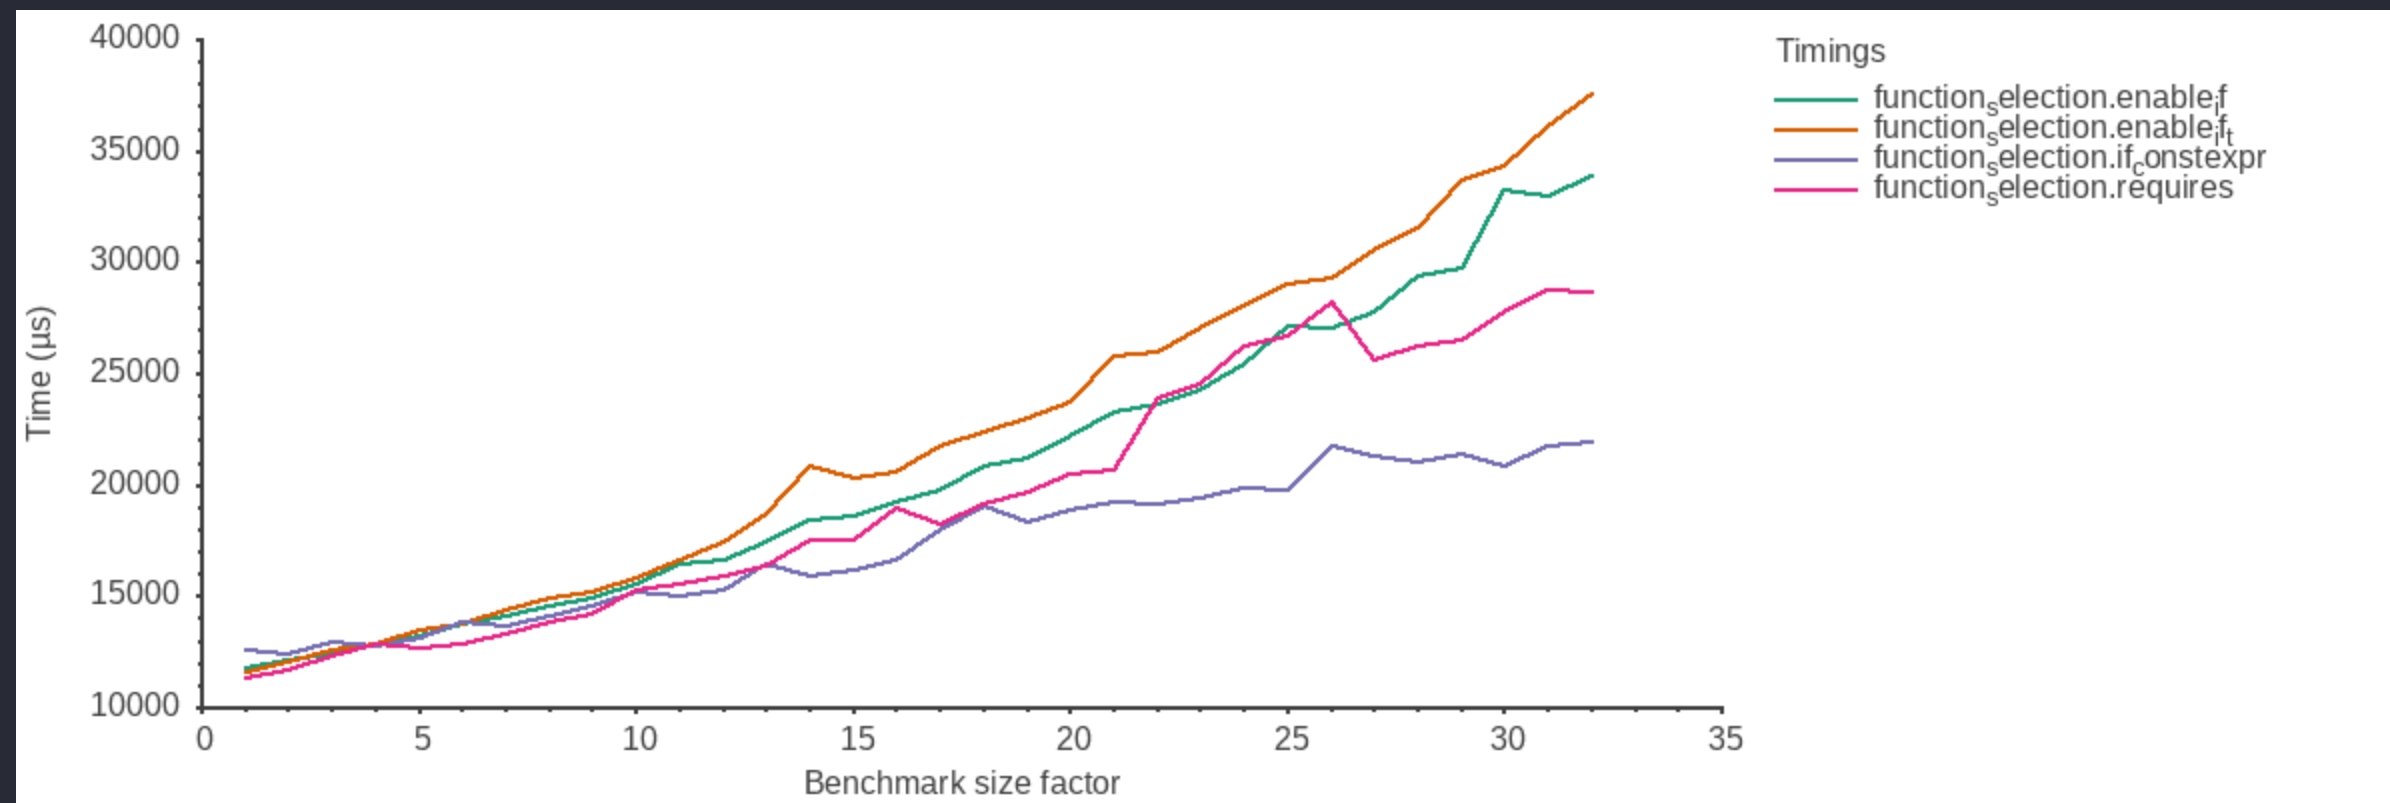
→ Based on library implementation comparisons in **Metabench**

ctbench

compile time benchmarking for Clang

C++ contenders - Function selection

enable_if_t vs enable_if vs if constexpr vs requires




Targeted data: ExecuteCompiler

ctbench

compile time benchmarking for Clang

C++ contenders - Function selection

enable_if_t vs enable_if vs if constexpr vs requires



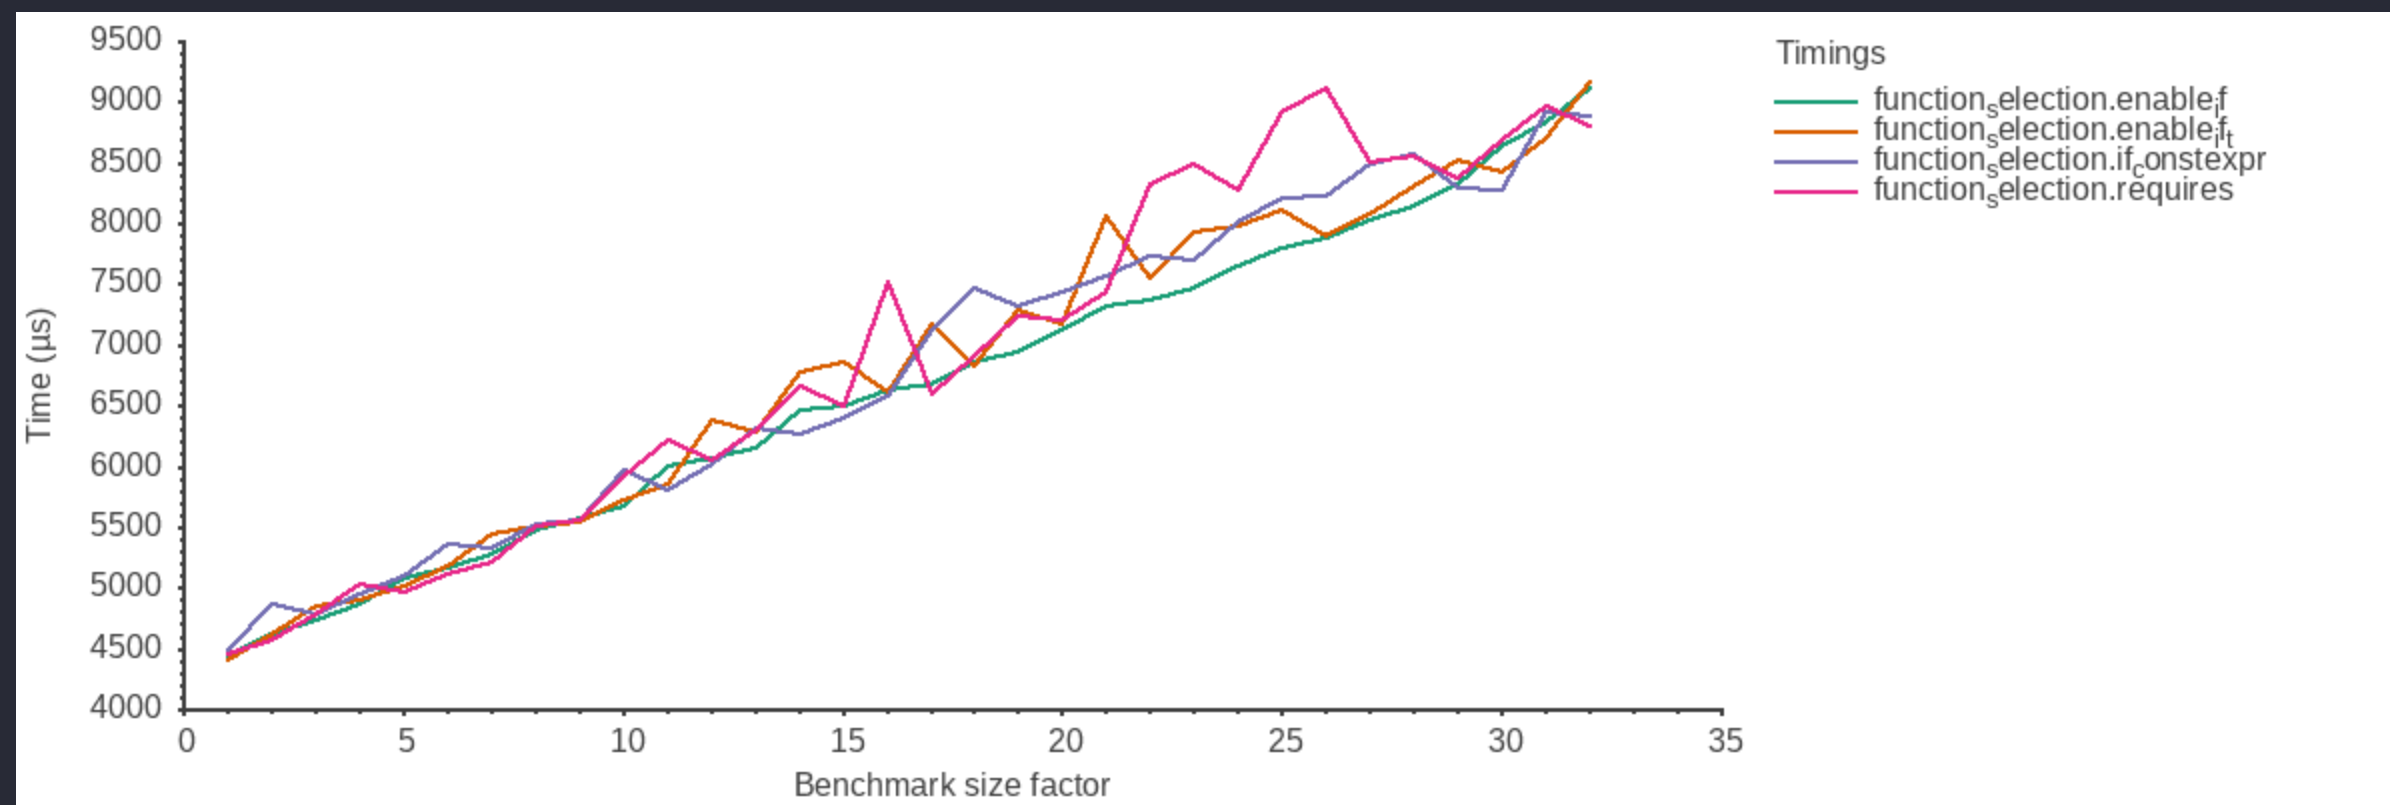
Targeted data: Frontend

ctbench

compile time benchmarking for Clang

C++ contenders - Function selection

enable_if_t vs enable_if vs if constexpr vs requires



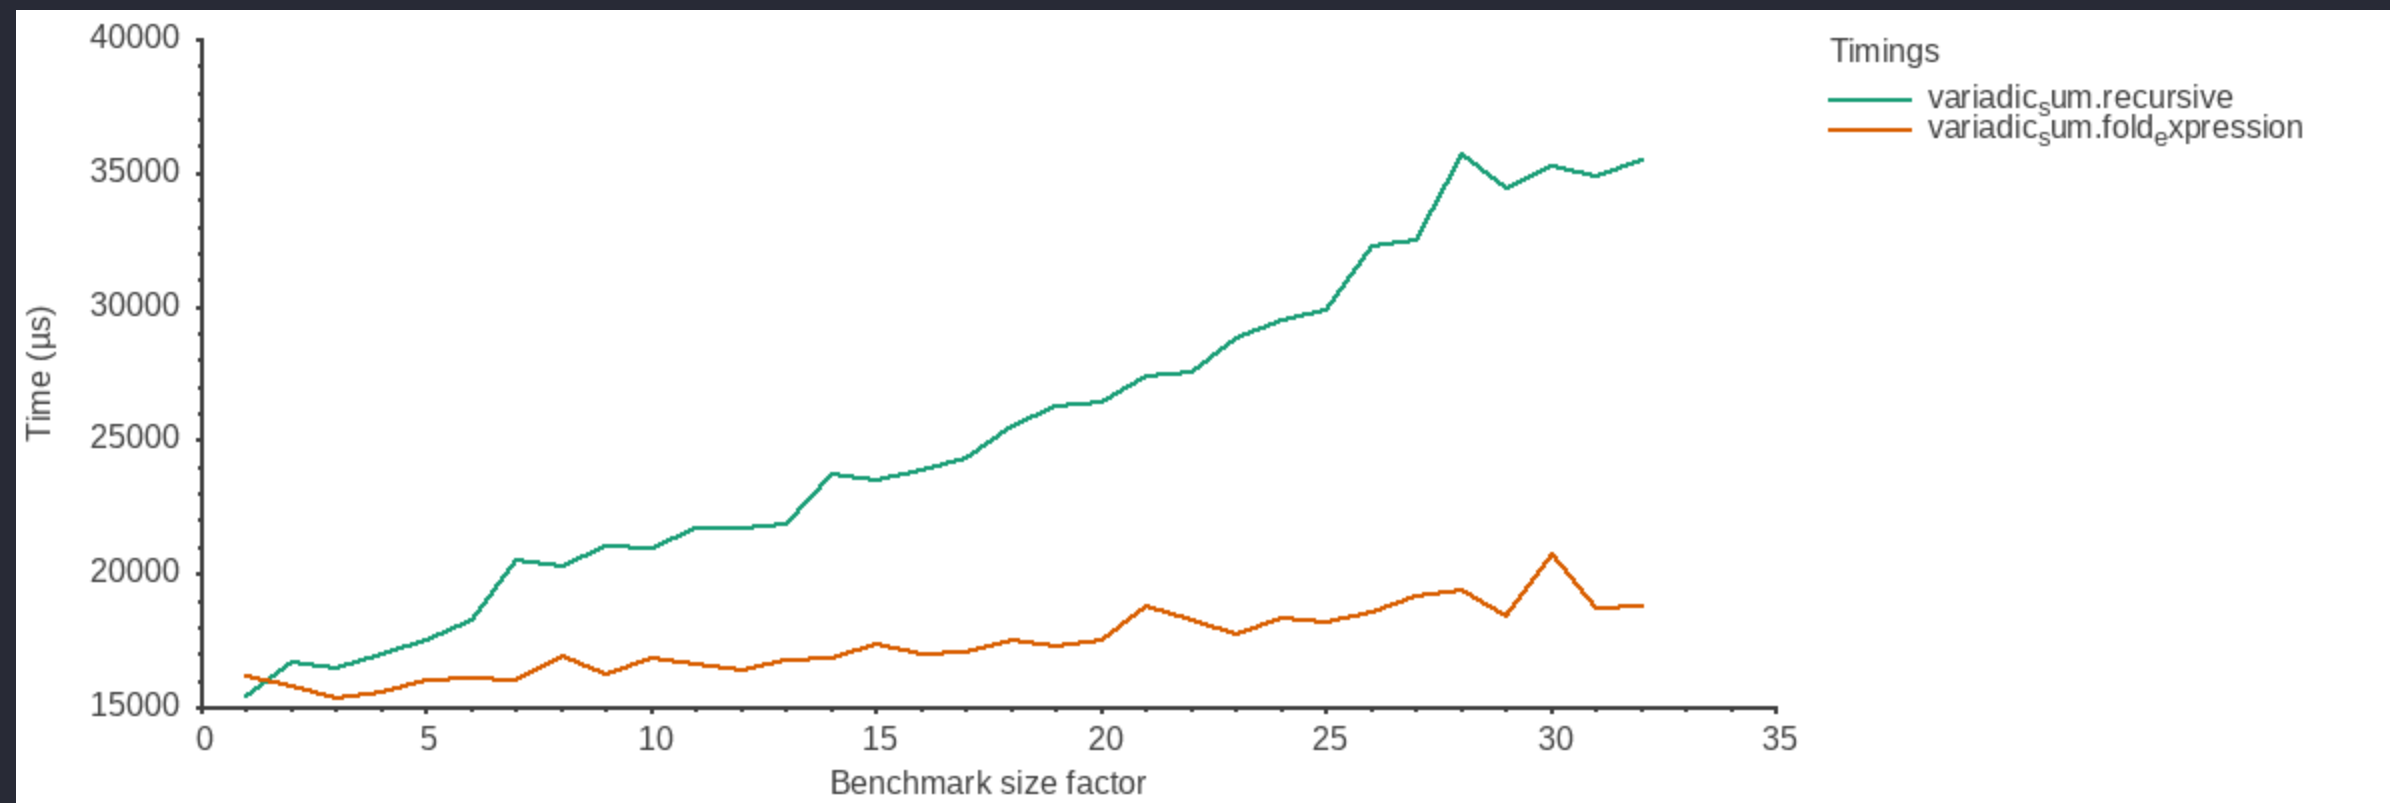
Targeted data: Backend

ctbench

compile time benchmarking for Clang

C++ contenders - Variadic sum

recursive vs fold_expression




Targeted data: ExecuteCompiler

ctbench

compile time benchmarking for Clang

C++ contenders - Variadic sum

recursive vs fold_expression



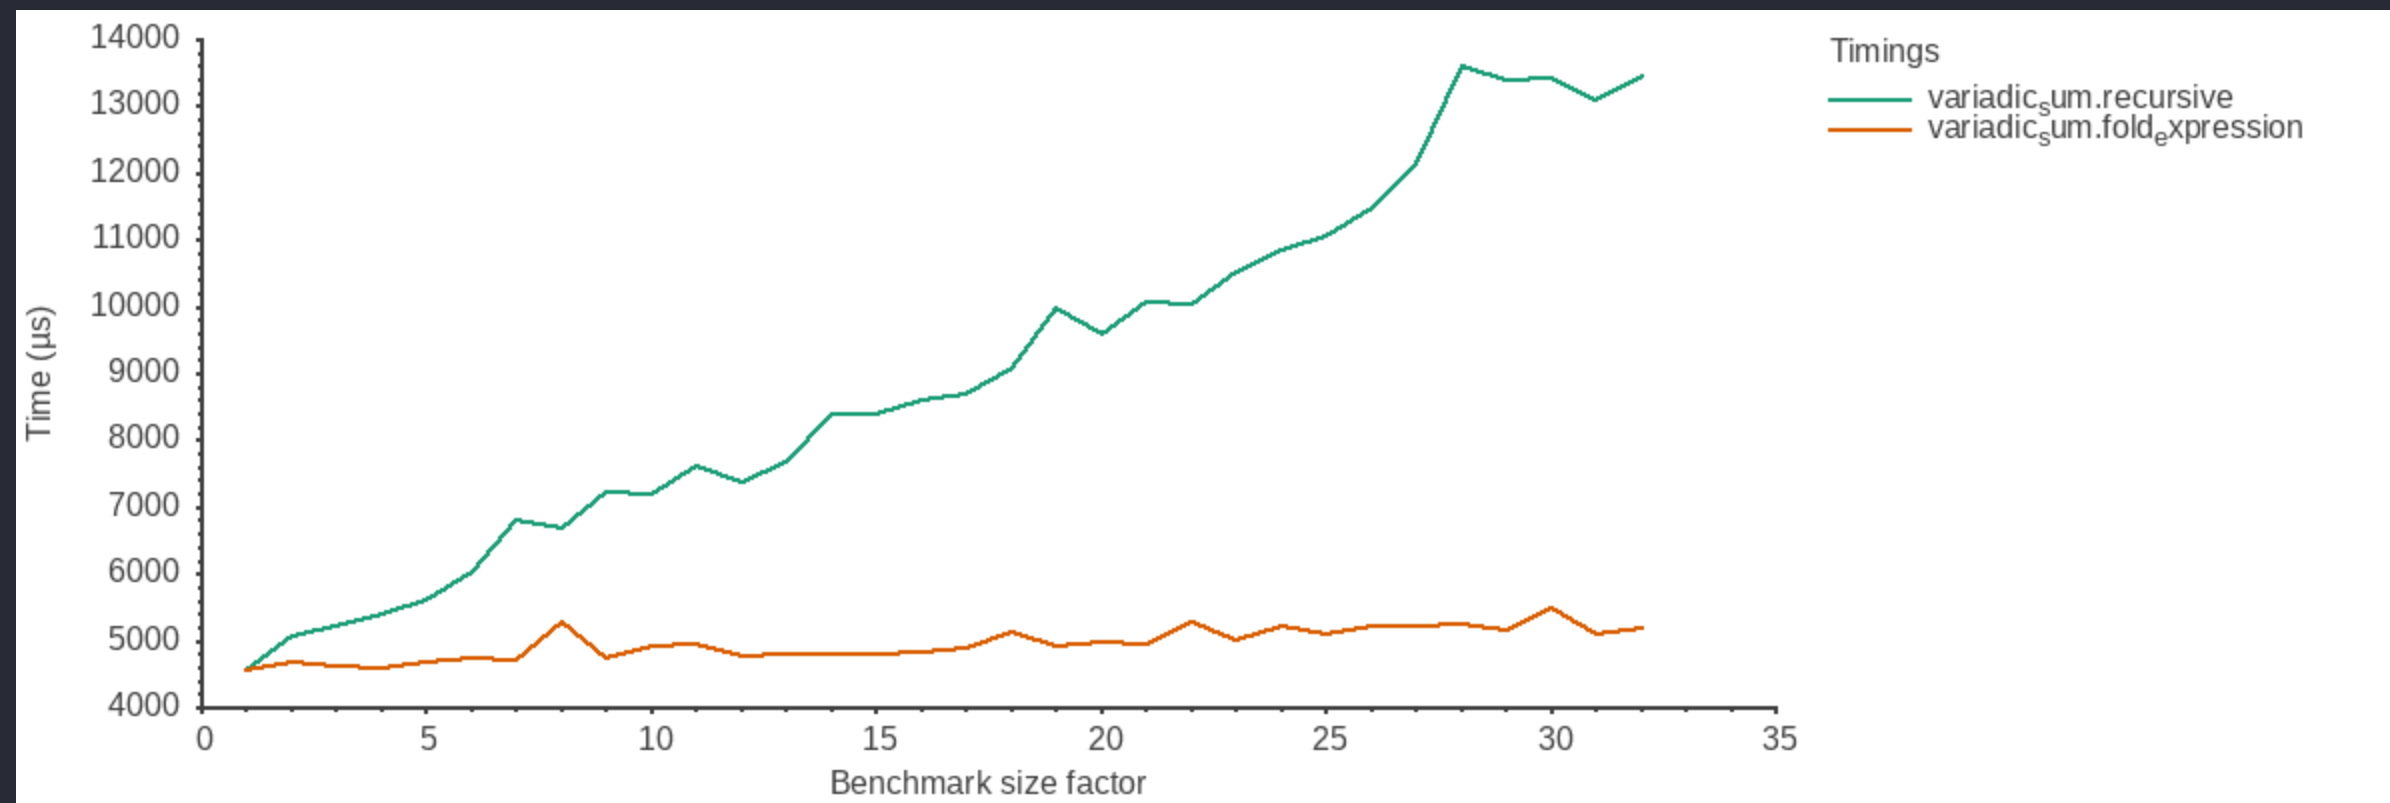
Targeted data: Frontend

ctbench

compile time benchmarking for Clang

C++ contenders - Variadic sum

recursive vs fold_expression



Targeted data: Backend

ctbench

compile time benchmarking for Clang

Project overview

- CMake API
 - **benchmarking.cmake** declares the end-user API,
 - documentation is provided inside (easily extracted into a MD file)
- **grapher** subproject (meatiest part)
 - CLI, time-trace file reading, predicate engine, and plotting,
 - designed as a **library** + CLI drivers,
 - relies heavily on **Sciplot** (<https://github.com/Sciplot/Sciplot>),
 - new plotters can be written easily
- Tooling:
 - **time-trace-wrapper**: clang exec wrapper to extract time-trace files
 - **cmake-doc-extractor**: extracts the API doc into a MD file

ctbench

compile time benchmarking for Clang

- What is the purpose of **ctbench**?
 - Better measurement and analysis tools where it's lacking, for a better **understanding** of compile times
- What do we have?
 - A flexible toolset to **run and analyze** Clang/LLVM compile-time benchmarks
- Is it perfect?
 - time-trace has **blindspots**, and we can't ignore **GCC**

ctbench

compile time benchmarking for Clang

Thank you!

<https://github.com/jpenuchot/ctbench>